



# *Java Database Connectivity...*

- Programs written in Java are able to communicate with relational databases (whether local or remote) via the Java Database Connectivity (JDBC) API

- In order to use JDBC for the accessing of data from a particular type of relational database, it is necessary to provide some mediating software that will allow JDBC to communicate with the vendor-specific API for that database
- Such software is referred to as a ***driver***

- Suitable drivers are usually supplied either by the database vendors themselves or by third parties
- ODBC drivers were originally available only for Microsoft (MS) databases, other vendors and third party suppliers have since brought out ODBC drivers for most of the major non-MS databases

- The most commonly used version of JDBC is currently JDBC 2, though there are still plenty of JDBC 1 drivers around, as well as an increasing number of JDBC 3 drivers

# Using JDBC to access a database requires several steps

1. Load the database driver.
2. Establish a connection to the database.
3. Use the connection to create a *Statement* object and store a reference to this object

4. Use the above *Statement* reference to run a specific query or update statement and accept the result(s)
5. Manipulate and display the results (if a query) or check/show number of database rows affected (for an update)
6. Repeat steps 4 and 5 as many times as required for further queries/updates.
7. Close the connection.

## *Load the Database Driver*

- This is achieved via static method ***forName*** of class *Class*

```
Class.forName("com.mysql.jdbc.Driver")
```



## ***Establish a Connection to the Database***

- We declare a ***Connection*** reference and call static method ***getConnection*** of class ***DriverManager*** to return a ***Connection*** object for this reference. Method ***getConnection*** takes three *String* arguments:
  - a URL-style address for the database
  - a user name;
  - a password.

Example:

Connection con =

```
DriverManager.getConnection("jdbc:mysql://  
localhost/database1","root", "");
```

## Create a Statement Object and Store its Reference

- A *Statement* object is created by calling the ***createStatement*** method of our *Connection* object
- The address of the object returned by this call to *createStatement* is saved in a *Statement* reference



Example:

Statement stm = con.createStatement ( );



## *Run a Query/Update and Accept the Result(s)*

- DML (Data Manipulation Language) statements in SQL may be divided into two categories:
- Those that retrieve data from a database (i.e., SELECT statements)
- and
- Those that change the contents of the database in some way (INSERT, DELETE and UPDATE statements)

- Class *Statement* has methods ***executeQuery*** and ***executeUpdate*** that are used to execute these two categories respectively
- The former method returns a ***ResultSet*** object, while the latter returns an integer that indicates the number of database rows that have been affected by the updating operation

- It is common practice to store the SQL query in a ***String*** variable and then invoke ***executeQuery*** with this string as an argument, in order to avoid a rather cumbersome invocation line

Example 1:

```
String selectAll = "SELECT * FROM Accounts";  
ResultSet results =stm.executeQuery(selectAll);
```

Example 2:

```
String selectFields = "SELECT acctNum, balance FROM  
    Accounts";  
ResultSet results = stm.executeQuery(selectFields);
```



## ***Manipulate/Display/Check Result(s)***

- The *ResultSet* object returned in response to a call of *executeQuery* contains the database rows that satisfy the query's search criteria
- Having moved to the particular row of interest via any of the methods, we can retrieve data via either the field name or the field position

- `int getInt (String <columnName>)`
- `int getInt (int <columnIndex>)`
- `String getString (String <columnName>)`
- `String getString (int <columnIndex>)`

- Initially, the *ResultSet* cursor/pointer is positioned **before** the first row of the query results, so method *next* must be called before attempting to access the results
- Such rows are commonly processed via a while loop that checks the Boolean return value of this method first (to determine whether there is any data at the selected position)



**Example Table...**



```
String select = "SELECT * FROM Accounts";
ResultSet results = stm.executeQuery(select);
while (results.next())
{
    System.out.println("Account no." + results.getInt(1));

    System.out.println("Account holder: " + results.getString(3) +
        " " + results.getString(2));

    System.out.println("Balance: " + results.getFloat(4));

    System.out.println ();
}
```

**NOTE: Column/field numbers start at 1, not 0**

- Alternatively, column/field names can be used

For example:

```
System.out.println("Account no." +  
results.getInt("acctNum");
```

## ***Close the Connection***

- This is achieved by calling method *close* of our *Connection* object and should be carried out as soon as the processing of the database has finished

```
con.close();
```

- Statement objects may also be closed explicitly via the identically-named method of our *Statement* object.

For example:

```
stm.close();
```



# Modifying the Database Contents

- INSERT, DELETE and UPDATE statements
- We shall have to submit our SQL statements via the *executeUpdate* method

### Example 1:

```
String s1 = "INSERT INTO Accounts" + " VALUES (123456,'Smith'," + "'John  
James',752.85)";
```

```
int result = stmt.executeUpdate(s1);
```

### Example 2:

```
String s2 = "UPDATE Accounts" + " SET surname = 'Bloggs'," + "firstNames  
= 'Fred Joseph'" + " WHERE acctNum = 123456";
```

```
stmt.executeUpdate(s2);
```

Example 3:

```
String s3 = "DELETE FROM Accounts"  
+ " WHERE balance < 100";
```

```
result = stmt.executeUpdate(s3);
```

```
int result = statement.executeUpdate(insert);  
if (result==0)  
    System.out.println("* Insertion failed! *");
```

# Transactions...

- A transaction is one or more SQL statements that may be grouped together as a single processing entity
- If only some of the statements are executed, then the database is likely to be left in an inconsistent state

- The SQL statements used to implement transaction processing are **COMMIT** and **ROLLBACK**, which are mirrored in Java by the *Connection* interface methods ***commit*** and ***rollback***

- ***commit*** is used at the end of a transaction to commit/finalize the database changes
- ***rollback*** is used (in an error situation) to restore the database to the state it was in prior to the current transaction (by undoing any statements that may have been executed)

- By default, however, JDBC automatically commits each individual SQL statement that is applied to a database
- In order to change this default behavior so that transaction processing may be carried out, we must first execute **Connection** method **setAutoCommit** with an argument of false (to switch off auto-commit).



```
.....  
con.setAutoCommit(false);  
  
.....  
try  
{  
    //Assumes existence of 3 SQL update strings  
    //called update1, update2 and update3.  
    stmt.executeUpdate(update1);  
    stmt.executeUpdate(update2);  
    stmt.executeUpdate(update3);  
    con.commit();  
}  
catch(SQLException sqlEx)  
{  
    con.rollback();  
    System.out.println("* SQL error! Changes aborted... *");  
}  
  
.....
```

## Meta Data...

- Meta data is 'data about data'
- There are two categories of meta data available through the JDBC API:
  - Data about the rows and columns returned by a query (i.e., data about *ResultSet* objects);
  - Data about the database as a whole.

- The first of these is provided by interface ***ResultSetMetaData***, an object of which is returned by the *ResultSet* method ***getMetaData***

- Information available from a *ResultSetMetaData* object includes the following:
- the number of fields/columns in a *ResultSet* object;
- the name of a specified field;
- the data type of a field;
- the maximum width of a field;
- the table to which a field belongs.

- Data about the database as a whole is provided by interface ***DatabaseMetaData***, an object of which is returned by the *Connection* method ***getMetaData***
- However, most Java developers will rarely find a need for ***DatabaseMetaData***

- `int getColumnCount()`
- `String getColumnName(<colNumber>)`
- `String getColumnName(<colNumber>)`

## Scrollable *ResultSet*s in JDBC 2...

- In all our examples so far, movement through a *ResultSet* object has been confined to the forward direction only, and even that has been restricted to moving by one row at a time
- With the emergence of JDBC 2 in Java 2, however, a great deal more flexibility was made available to Java programmers by the introduction of the following *ResultSet* methods:



`boolean first()`

`boolean last()`

`boolean previous()`

`boolean relative (int <rows>)`

`boolean absolute(int <rows>)`





For the first 3 methods, as with method *next*, the return value in each case indicates whether or not there is data at the specified position

- Method ***relative*** takes a signed argument and moves forwards/backwards the specified number of rows

For example:

```
results.relative(-3);    //Move back 3 rows.
```

- Method ***absolute*** also takes a signed argument and moves to the specified absolute position, counting either from the start of the *ResultSet* (for a positive argument) or from the end of the *ResultSet* (for a negative argument)

*For example:*

```
results.absolute(3);
```

```
//Move to row 3 (from start of ResultSet).
```

- Before any of these new methods can be employed, however, it is necessary to create a *scrollable ResultSet*
- This is achieved by using an overloaded form of the Connection method *createStatement* that takes two integer arguments

Statement *createStatement*(int <resultSetType>, int  
<resultSetConcurrency>)

- There are three possible values that the first argument can take to specify the type of *ResultSet* object that is to be created
- These three values are identified by the following static constants in interface *ResultSet*:

TYPE\_FORWARD\_ONLY

TYPE\_SCROLL\_INSENSITIVE

TYPE\_SCROLL\_SENSITIVE

- The first option allows only forward movement through the *ResultSet*
- The second and third options allow movement of the *ResultSet's cursor* both forwards and backwards through the rows
- The difference between these two is that *TYPE\_SCROLL\_SENSITIVE* causes any changes made to the data rows to be reflected dynamically in the *ResultSet* object, while *TYPE\_SCROLL\_INSENSITIVE* does not

- There are two possible values that the second argument to *createStatement* can take
- These are identified by the following static constants in interface *ResultSet*:  

*CONCUR\_READ\_ONLY*  
*CONCUR\_UPDATABLE*
- The first means that we cannot make changes to the *ResultSet* rows, while the second will allow changes to be made

## Modifying Databases via Java Methods...

***Another very useful feature of JDBC 2 is the ability to modify ResultSet rows directly via Java methods (rather than having to send SQL statements), and to have those changes reflected in the database itself***



- In order to do this, it is necessary to use the second version of *createStatement* again (i.e., the version that takes two integer arguments) and supply *ResultSet.CONCUR\_UPDATABLE* as the second argument
- The updateable *ResultSet* object does not have to be scrollable, but, when making changes to a *ResultSet*, we often want to move freely around the *ResultSet* rows, so it seems sensible to make the *ResultSet* scrollable

## Example

```
Statement stm = con.createStatement(  
    ResultSet.TYPE_SCROLL_SENSITIVE,  
    ResultSet.CONCUR_UPDATABLE);
```

As usual, there are three types of change that we can carry out on the data in a database:

- updates (of some/all fields of a selected row)
- insertions (of new data rows)
- deletions (of existing database rows)

- There is a set of ***update\*\*\**** methods (analogous to the *get\*\*\** methods that we use to retrieve the data from a row within a *ResultSet*), each of these methods corresponding to one of the data types that may be held in the database
- For example, there are methods ***updateString*** and ***updateInt*** to update *String* and *int* data respectively

- Each of these methods takes two arguments:
  - A string specifying the name of the field to be updated;
  - A value of the appropriate type that is to be assigned to the field

- There are three steps involved in the process of **updating**:
  - position the **ResultSet** *cursor at the required row*;
  - call the appropriate **update\*\*\*** *method(s)*;
  - call method **updateRow**

```
results.absolute(2); //Move to row 2 of ResultSet  
results.updateFloat("balance", 42.55f);  
results.updateRow();
```

Note here that an 'f' must be appended to the float literal, in order to prevent the compiler from interpreting the value as a double

- For an **insertion**, the new row is initially stored within a special buffer called the 'insertion row' and there are three steps involved in the process:
  - call method *moveToInsertRow*;
  - call the appropriate **update**\*\*\* *method for each field in the row*;
  - call method *insertRow*.

```
results.moveToInsertRow();  
results.updateInt("acctNum", 999999);  
results.updateString("surname", "Harrison");  
results.updateString("firstNames", "Christine Dawn");  
results.updateFloat("balance", 2500f);  
results.insertRow();
```

- **Get\*\*\*** methods called after insertion will not retrieve values for newly-inserted rows
- If this is the case with a particular database, then it will be necessary to close the **ResultSet** and create a new one (using the original query), in order for the new insertions to be recognized



- To delete a row without using SQL, there are just two steps:
  - move to the appropriate row
  - call method ***deleteRow***

Example:

```
results.absolute(3);           //Move to row 3.  
results.deleteRow();
```

- Note that JDBC drivers can handle deletions differently
- Some remove the row completely from the *ResultSet*, while others use a blank row as a placeholder
- *With* the latter, the original row numbers are not changed



*Thank You . . .*

